

NLP Assignment-1

Dr. Dheeraj

March 6, 2025

(Roll Nos: 1-5)

1. Develop an NLP pipeline that performs:

- a) Text preprocessing.
- b) Given a word and its possible meanings, select the correct sense based on a sentence.
- c) Compare the overlapping words in the sentence and the sense definition.
- d) Read the research papers.

Text Preprocessing

- ▶ **Input:** "The quick brown fox jumps over the lazy dog!"
- ▶ **Steps (Examples):**
 1. **Tokenization:** ['The', 'quick', 'brown', 'fox', 'jumps', 'over', 'the', 'lazy', 'dog', '!']
 2. **Lowercasing:** ['the', 'quick', 'brown', 'fox', 'jumps', 'over', 'the', 'lazy', 'dog', '!']
 3. **Stopword Removal:** ['quick', 'brown', 'fox', 'jumps', 'lazy', 'dog', '!']
 4. **Lemmatization:** ['quick', 'brown', 'fox', 'jump', 'lazy', 'dog']
- ▶ **Output:** ['quick', 'brown', 'fox', 'jump', 'lazy', 'dog']

Word Sense Disambiguation

- ▶ **Input Sentence:** "He sat on the bank of the river."
- ▶ **Word to Disambiguate:** "bank"
- ▶ **Possible Meanings:**
 1. **bank (financial institution):** A financial establishment.
 2. **bank (river edge):** The slope beside a body of water.
- ▶ **Context Words:** "sat", "river"
- ▶ **Selected Sense:** bank (river edge)

Overlapping Words Comparison

- ▶ **Input Sentence:** "He sat on the bank of the river."
- ▶ **Word Definition:** "The slope beside a body of water."
- ▶ **Extracted Words from Sentence:** ["sat", "bank", "river"]
- ▶ **Extracted Words from Definition:** ["slope", "body", "water"]
- ▶ **Common Words:** "river" matches with "water" (semantic similarity)
- ▶ **Number of Overlapping Words:** 1
- ▶ **Similarity Score:** 0.3 (depending on metric used)

(Roll Nos: 6-10)

2. Develop an NLP pipeline that performs:

- ▶ **(c)** Generate Text N-Grams Without Using NLTK:
 - ▶ Implement n-gram extraction (bi-grams, tri-grams, etc.) from a given text.
- ▶ **(d)** Ignore stopwords and punctuation in the process.

Task (c): Generate Text N-Grams

Input:

- ▶ A given long text (e.g., "The quick brown fox jumps over the lazy dog.").
- ▶ The value of **N** (e.g., 2 for bigrams, 3 for trigrams).

Process:

- ▶ Tokenize the text manually (split into words).
- ▶ Ignore punctuation.
- ▶ Generate n-grams by forming consecutive sequences of words of size N.

Output:

- ▶ List of n-grams.
- ▶ **Example (bigrams):** [("The", "quick"), ("quick", "brown"), ("brown", "fox"), ...]
- ▶ **Example (trigrams):** [("The", "quick", "brown"), ("quick", "brown", "fox"), ...]

Task (d): Ignore Stopwords and Punctuation

Input:

- ▶ The same input text.
- ▶ List of stopwords (e.g., ["the", "is", "and", "over", "in", "on"]).

Process:

- ▶ Remove stopwords and punctuation before generating n-grams.

Output:

- ▶ N-grams **without stopwords and punctuation**.
- ▶ **Example Before Stopword Removal:** [("The", "quick"), ("quick", "brown"), ("brown", "fox")]
- ▶ **Example After Stopword Removal:** [("quick", "brown"), ("brown", "fox")]

(Roll Nos: 11-15)

Task 3: Develop an NLP pipeline that performs:

1. Compute TF-IDF scores for each word in a document.
2. Return the top N keywords with the highest scores.

Step (c): Compute TF-IDF Scores

Input:

- ▶ A document (text data).
- ▶ Preprocessed text (tokenized, lowercased, stopwords removed).

Processing:

- ▶ Compute Term Frequency (TF):

$$TF = \frac{\text{Count of word in document}}{\text{Total words in document}}$$

- ▶ Compute Inverse Document Frequency (IDF):

$$IDF = \log \left(\frac{\text{Total number of documents}}{\text{Number of documents containing the word}} \right)$$

- ▶ Calculate TF-IDF:

$$TF - IDF = TF \times IDF$$

Output: List of words with their TF-IDF scores.

Step (d): Extract Top N Keywords

Input:

- ▶ List of words with TF-IDF scores.
- ▶ Value of N (number of top keywords).

Processing:

- ▶ Sort words by TF-IDF scores in descending order.
- ▶ Select the top N highest-scoring words.

Output: A ranked list of top N keywords.

Example: TF-IDF Calculation

Input Document:

"Natural Language Processing (NLP) is a field of AI that focuses on the interaction between humans and computers using natural language. NLP techniques include tokenization, parsing, and named entity recognition."

Output (Top 3 Keywords)

Keyword	TF-IDF Score
NLP	0.89
Natural Language	0.75
Processing	0.68

4. Develop an NLP pipeline that performs:

- ▶ (c) Implement extractive text summarization by scoring sentences based on:
 - ▶ Word frequency
 - ▶ Sentence length
 - ▶ TF-IDF scores
- ▶ (d) Find the most frequent POS (Part of Speech) tag in a given text.

(c) Extractive Text Summarization - Hint

Input:

- ▶ A large text document (e.g., an article, research paper, or long paragraph).

Processing:

- ▶ Tokenization (split text into sentences and words).
- ▶ Compute word frequency, sentence length, and TF-IDF scores.
- ▶ Rank sentences based on scores.
- ▶ Select the top-ranked sentences for the summary.

Output:

- ▶ A concise summary containing the most important sentences.

Example: Extractive Summarization

Input Text: *"Natural Language Processing (NLP) is a field of AI that focuses on the interaction between humans and computers using natural language. It enables applications such as text summarization, machine translation, and chatbots."*

Extracted Summary: *"NLP is a field of AI focusing on human-computer interaction. It enables text summarization, translation, and chatbots."*

(d) Most Frequent POS Tag - Hint

Input:

- ▶ A given text document or paragraph.

Processing:

- ▶ Tokenize text into words.
- ▶ Perform POS tagging using NLP libraries.
- ▶ Count occurrences of each POS tag.
- ▶ Identify the most frequent POS tag.

Output:

- ▶ The most frequent POS tag in the text.

Example: Most Frequent POS Tag

Input Text: *"The cat sat on the mat and looked at the window."*

POS Tags Count:

- ▶ Nouns (NN): 3
- ▶ Verbs (VB): 2
- ▶ Determiners (DT): 2

Output: Most frequent POS tag: **Noun (NN)**

(Roll Nos: 21-25)

5. Develop an NLP pipeline that performs:

1. **(c)** Apply **Byte Pair Encoding (BPE)** to compress text data and reduce vocabulary size.
2. **(d)** Given a BPE-encoded text, reconstruct the original words by merging subwords.

Byte Pair Encoding (BPE) for Text Compression

Step (c): Apply BPE to Reduce Vocabulary Size

- ▶ **Input:** A large text corpus (e.g., thousands of sentences).
- ▶ **Hyperparameter - Merge Operations:**
 - ▶ Defines how many times frequent subwords are merged.
 - ▶ Higher merges \Rightarrow Larger subwords, smaller vocabulary.
 - ▶ Lower merges \Rightarrow Smaller subwords, larger vocabulary.
- ▶ **Processing:**
 1. Tokenize text into individual characters or subwords.
 2. Identify and merge the most frequent adjacent subwords **N** times.
 3. Store the learned subword vocabulary.
- ▶ **Output:**
 - ▶ Compressed text with subword units.
 - ▶ Vocabulary of subwords.

Example of BPE Encoding

Original Sentence: "learning NLP is useful"

After 2 BPE Merge Operations:

▶ "learn@@ ing NLP is use@@ ful"

Reconstructing Original Words from BPE

Step (d): Decode BPE-encoded text

- ▶ **Input:** BPE-encoded text and learned merge rules.
- ▶ **Processing:**
 1. Reverse the BPE process.
 2. Merge subwords iteratively using learned rules.
- ▶ **Output:** Original words restored.

Example:

- ▶ **Input:** "learn@@ ing NLP is use@@ ful"
- ▶ **Output:** "learning NLP is useful"

6. (Roll Nos: 26-33)

6. Develop an NLP pipeline that performs:

- (c) Find the most frequently occurring word in a text, excluding common stopwords.
- (d) Find the most semantically similar sentence to a given input using TF-IDF.

Task (c): Most Frequent Word (Excluding Stopwords)

Input:

- ▶ A text document (long paragraph or multiple sentences).
- ▶ A list of stopwords (e.g., from NLTK or SpaCy).

Processing Steps:

- ▶ Tokenize text into words.
- ▶ Convert words to lowercase and remove punctuation.
- ▶ Remove stopwords.
- ▶ Count word occurrences and find the most frequent word.

Output: The most frequently occurring word (excluding stopwords).

Example for Task (c)

Input: *"The cat sat on the mat. The cat is happy. The dog is happy too."*

After Removing Stopwords: *"cat sat mat. cat happy. dog happy."*

Output: **"cat"** (Most frequently occurring word)

Task (d): Sentence Similarity using TF-IDF

Input:

- ▶ A collection of sentences (text corpus).
- ▶ A query sentence to compare.

Processing Steps:

- ▶ Compute TF-IDF vectors for all sentences in the corpus.
- ▶ Compute TF-IDF vector for the query sentence.
- ▶ Calculate cosine similarity between query and each sentence.
- ▶ Select the most similar sentence.

Output: The sentence with the highest similarity score.

Example for Task (d)

Corpus Sentences:

1. "The sky is blue and beautiful."
2. "The weather is amazing today."
3. "The sun is shining brightly in the sky."

Query Sentence: *"It is a bright sunny day."*

Most Similar Sentence: *"The sun is shining brightly in the sky."*

Summary

Key Points:

- ▶ Task (c): Identify the most frequently occurring word after removing stopwords.
- ▶ Task (d): Find the sentence most similar to a query using TF-IDF and cosine similarity.